

STK4900 Oblig1

Daniel Heinesen, daniehei

March 9, 2019

Exercise 1

a)

	log.no2	log.cars	temp	wind.speed	hour.of.day
1	Min. :1.224	Min. :4.127	Min. :-18.6000	Min. :0.300	Min. : 1.00
2	1st Qu.:3.214	1st Qu.:6.176	1st Qu.: -3.9000	1st Qu.:1.675	1st Qu.: 6.00
3	Median :3.848	Median :7.425	Median : 1.1000	Median :2.800	Median :12.50
4	Mean :3.698	Mean :6.973	Mean : 0.8474	Mean :3.056	Mean :12.38
5	3rd Qu.:4.217	3rd Qu.:7.793	3rd Qu.: 4.9000	3rd Qu.:4.200	3rd Qu.:18.00
6	Max. :6.395	Max. :8.349	Max. : 21.1000	Max. :9.900	Max. :24.00

Table 1

Statistic	N	Mean	St. Dev.	Min	Pctl(25)	Pctl(75)	Max
log.no2	500	3.698	0.751	1.224	3.214	4.217	6.395
log.cars	500	6.973	1.087	4.127	6.176	7.793	8.349
temp	500	0.847	6.525	-18.600	-3.900	4.900	21.100
wind.speed	500	3.056	1.784	0.300	1.675	4.200	9.900
hour.of.day	500	12.382	6.803	1	6	18	24

Exercise 2

a)

The data we are looking at is the systolic blood pressure for 36 people divided into three groups determined by their age. These groups are 30 – 45 years, 46 – 59 years and 60 – 75 years.

Age Group	N	Mean	St. Dev.	Min	Pctl(25)	Pctl(75)	Max
All	36	138.806	25.749	104	117.5	156.2	214
30 – 45 years	12	122.167	15.338	104	112	129	160
46 – 59 years	12	139.083	22.625	108	121.5	157.8	174
60 – 75 years	12	155.167	27.719	110	138	164	214

Table 2: A table showing different measures for the distribution of the systolic blood pressure for the three age groups, and the entire group as a whole. The increase in all the measures for the higher age groups seems to indicate a correlation between age and blood pressure.

If we look at table 1 we see the numerical summary of the blood pressure of the three age groups, and the group as a whole. From the mean column, we see that mean blood pressure seems to increase with increasing age group. This is also the case for all the other measures see in the table. This seems to indicate that there is some relationship between systolic blood pressure and age.

This may be just repetition of the caption...




Figure 1: Boxplots showing the distribution of the blood pressure of the three age groups.

From fig. 1 we see that there is a clear increase in blood pressure with age. Both the median, max and min increases with age. This is exactly what we saw in tab. 1 – a difference is that in the table the mean was given, while in the box plot the median is given. What we can see clearer from the plot is that the distributions also become wider as age increase. The interquartile width of the middle age group is quite wide, and overlaps with the other groups. This can make this group a bit more difficult to get a significant result from.

All in all, both the numerical summary, tab. 1, and the box plot, fig. 1, seems to show the same: A increase in systolic blood pressure as age increases.

b)

We now want to look close at whether blood pressure varies between the different age groups. Above we used qualitative reasoning to to argue that this is the case. Now we want a quantified measure of this. To do this we will use a one-way ANOVA.

For the ANOVA we have the following hypotheses:

- $H_0 : \mu_1 = \mu_2 = \mu_3$
- $H_a : \mu_1 \neq \mu_2 \neq \mu_3$

In word: We have a null hypothesis that the mean blood pressure of all the groups are the same, and an alternative hypothesis that this is not that case. For the ANOVA to be valid, we need to assume that:

- The observations are independent of one another
- The observations in one group are a random sample with a normal distribution $N(\mu_k, \sigma^2)$

From tab. 2 we see that we get a small p-value of $p = 0.0043$. This is below the $p = 0.05$ mark which is standard to use as the limit of significance. This means that we can throw out that null hypothesis, and conclude that blood pressure indeed varies across the age groups ¹

¹It is normally bad practice to interpret the p-value in an article, but since this is an oblig, it is done here.

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Age Group	2	6535.39	3267.69	6.47	0.0043
Residuals	33	16670.25	505.16		

Table 3: A table showing the result of an ANOVA test on the blood pressure dataset. We see that we have a high F-value, leading to a small, significant p-value.

c)

We are now going to try to use a categorical regression on the data set. We are going to use treatment-contrast with group 1, the youngest, as the reference group. This means that our model will look like

$$y_i = \mu_1 + (\mu_2 - \mu_1) \cdot x_{1,i} + (\mu_3 - \mu_1) \cdot x_{2,i} + \epsilon_i, \quad (1)$$

where ϵ_i is a normally distributed error term, μ_j is the mean of the different age groups, and

$$x_{j-1,i} = \begin{cases} 1 & \text{for } i \text{ in group } j \\ 0 & \text{else} \end{cases}. \quad (2)$$

This means that if patient is in group 2, then $x_{1,i} = 1$ and $x_{2,i} = 0$, and vis versa if the patient is in group 3. If the patient is in group 1, both will be 0.

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	122.1667	6.4882	18.83	0.0000
Age Group 2	16.9167	9.1757	1.84	0.0742
Age Group 3	33.0000	9.1757	3.60	0.0010
Residual standard error: 22.48 on 33 degrees of freedom				
Multiple R-squared: 0.2816, Adjusted R-squared: 0.2381				
F-statistic: 6.469 on 2 and 33 DF, p-value: 0.004263				

Table 4: The summary of the categorical regression of the blood pressure with respect to the age groups.

In tab. 3 we see the summary of the regression model with the age groups as a categorical predictor variable. The intercept in the table is the mean of the youngest age group $\mu_1 = 122.167$ – our reference group –. The slopes of the two variables can be interpreted as how much the mean of blood pressure increases if you are in that age group, compared to the reference group. E.g. if you are in age group 2, we will expect your blood pressure to be 16.9 mmHg higher than the mean of the youngest group, i.e. it is expected to be 139.067 mmHg.

We can look at the p-values of tab. 3 we see that the p-value for age group 3 is $p = 0.001$. This means that the increase of 33 mmHg is significant. But for group 2 the p-value is just 0.0742, which means that this increase is not significant, so we can not say that we have an 17 mmHg increase compared with the reference. This might have to do with what we see in the box plots 1, where the distribution of the second age group is very wide. The reason might also be due to a non linear dependence on age.

We also see that we have a low $R^2 = 0.2816$, which means that the blood pressure is not fully described by the age groups.

But all in all we see that we have a good regression model for group 1 and 3, with the change in group 2 not being significant. So there is some significant variation between the age groups, as we saw with the ANOVA.

Look
over and
rewrte...